

Alphalens Documentation

Contents

1	Overview	3
2	alphalens.utils	3
2.1	Data model used by this module	3
2.2	Exceptions	3
2.3	Core workflow	3
2.4	Quantization and bucketing	5
2.5	Return and performance helpers	5
2.6	Calendar and time utilities	5
2.7	Column detection and display helpers	6
2.8	Error handling helper	7
2.9	Practical usage	7
2.10	Notes	7
3	alphalens.performance	7
3.1	Data model used by this module	8
3.2	Core analysis functions	8
3.3	Portfolio weighting and return simulation	8
3.4	Quantile and bucket analysis	11
3.5	Event-study helpers	13
3.6	Portfolio simulation outputs	14
3.7	Dependencies and implementation notes	15
3.8	Practical usage	15
4	alphalens.plotting	15
4.1	Styling and context management	15
4.2	Data table displays	16
4.3	Information Coefficient visualization	17
4.4	Returns analysis	18
4.5	Factor characteristics	20
4.6	Constants	20
4.7	Dependencies and conventions	21
4.8	Practical usage	21
4.9	Notes	21

5	alphalens.tears	21
5.1	Helper class	22
5.2	Tear sheet generators	22
5.3	Data model	25
5.4	Dependencies and concepts	26
5.5	Practical usage	26
5.6	Notes	26

1 Overview

Alphalens is organized as a pipeline of four core modules:

- `alphalens.utils` prepares factor data and computes forward returns.
- `alphalens.performance` analyzes factor behavior and simulated portfolio performance.
- `alphalens.plotting` visualizes metrics and diagnostics.
- `alphalens.tears` assembles the results into tear sheets.

The usual workflow is:

1. Build cleaned factor data with `alphalens.utils.get_clean_factor_and_forward_returns()`.
2. Compute diagnostics and portfolio statistics with `alphalens.performance`.
3. Visualize outputs with `alphalens.plotting`.
4. Generate presentation-ready reports with `alphalens.tears`.

All modules typically operate on a MultiIndex `pandas` object indexed by `(date, asset)`.

2 `alphalens.utils`

Utilities for preparing factor data, computing forward returns, binning factors into quantiles/bins, and working with custom trading calendars.

This module is the main data-preparation layer used throughout Alphalens. Most workflows start with `get_clean_factor_and_forward_returns()` and then feed the resulting MultiIndex `DataFrame` into the tear sheets and plotting functions.

2.1 Data model used by this module

Most functions operate on a `pandas` object indexed by:

- level 0: `date` / timestamp
- level 1: `asset`

Forward returns are usually stored in columns named with `pd.Timedelta`-compatible strings such as `1D`, `5D`, `30m`, or `1D1h`.

2.2 Exceptions

`NonMatchingTimezoneError` Raised when factor and pricing indices have different timezones.

`MaxLossExceededError` Raised by `get_clean_factor()` when the share of dropped factor data exceeds `max_loss`.

2.3 Core workflow

2.3.1 `compute_forward_returns(factor, prices, periods=(1, 5, 10), filter_zscore=None, cumulative_returns=True)`

Compute N-period forward returns for each asset in `factor`.

Input expectations

- **factor**: MultiIndex Series indexed by (date, asset).
- **prices**: wide DataFrame with dates on the index and assets on the columns.
- The price history must extend far enough beyond the factor timestamps to cover the largest requested period.

Behavior

- Validates that factor and price timezones match.
- Infers a trading calendar when the factor index has no frequency.
- Computes forward returns for each requested period.
- Optionally filters out extreme values using `filter_zscore`.
- Names the output columns using a Timedelta-style string.

Returns A MultiIndex DataFrame with the same (date, asset) index as the input factor.

2.3.2 `get_clean_factor(factor, forward_returns, groupby=None, binning_by_group=False, quantiles=5, bins=None, groupby_labels=None, max_loss=0.35, zero_aware=False)`

Align a factor series with already-computed forward returns and assign factor quantiles or bins.

Use this when you already have forward returns and do not need to recompute them.

Key behavior

- Joins factor values to forward returns.
- Optionally maps assets to groups.
- Computes factor quantiles or bins using `quantize_factor()`.
- Drops rows with missing data.
- Tracks how much data was lost and raises `MaxLossExceededError` if the loss exceeds `max_loss`.

Returns A cleaned MultiIndex DataFrame containing forward return columns, `factor`, optionally `group`, and `factor_quantile`.

2.3.3 `get_clean_factor_and_forward_returns(factor, prices, groupby=None, binning_by_group=False, quantiles=5, bins=None, periods=(1, 5, 10), filter_zscore=20, groupby_labels=None, max_loss=0.35, zero_aware=False, cumulative_returns=True)`

One-stop helper that computes forward returns and then cleans and bins the factor data.

This is the primary entry point for most Alphalens analyses.

Equivalent steps

1. Call `compute_forward_returns()`
2. Call `get_clean_factor()`

Returns A cleaned MultiIndex DataFrame suitable for Alphalens tear sheets and plots.

2.4 Quantization and bucketing

2.4.1 `quantize_factor(factor_data, quantiles=5, bins=None, by_group=False, no_raise=False, zero_aware=False)`

Assign factor values to quantile or value bins.

Behavior

- Exactly one of `quantiles` or `bins` must be provided.
- Supports group-wise binning with `by_group=True`.
- Supports zero-aware binning, which splits positive and negative values separately.
- Can suppress binning failures with `no_raise=True`, returning NaN for problematic rows.

Returns A Series named `factor_quantile` indexed by `(date, asset)`.

2.4.2 `non_unique_bin_edges_error(func)`

Decorator that adds a more helpful error message when `pandas` cannot create unique bin edges.

This is used internally around `quantize_factor()` to explain common binning failures, especially when many identical factor values span multiple quantiles.

2.5 Return and performance helpers

2.5.1 `backshift_returns_series(series, N)`

Shift a MultiIndex returns series backward by N observations in the first level.

This helper is useful for converting backward-looking returns into forward-looking returns.

2.5.2 `demean_forward_returns(factor_data, grouper=None)`

Demean forward returns by date or by a custom grouper.

Behavior

- If `grouper` is not provided, demeaning happens per date.
- Only forward-return columns are adjusted.
- The result preserves the original shape and index.

2.5.3 `rate_of_return(period_ret, base_period)`

Convert returns observed over one period length into an equivalent rate for `base_period`.

This is useful for normalizing returns across different horizons.

2.5.4 `std_conversion(period_std, base_period)`

Convert standard deviation or standard error from one period length to another.

2.6 Calendar and time utilities

2.6.1 `infer_trading_calendar(factor_idx, prices_idx)`

Infer a trading calendar from factor and price datetimes.

Behavior

- Detects which weekdays are traded.
- Infers holidays by comparing the observed timestamps to a weekday-specific business-day calendar.
- Returns a `CustomBusinessDay` offset.

2.6.2 `timedelta_to_string(timedelta)`

Convert a `pd.Timedelta` into a compact string compatible with `pd.Timedelta(...)`.

Example output formats include `1D`, `3h15m`, and `1D1h`.

2.6.3 `timedelta_strings_to_integers(sequence)`

Convert a sequence of timedelta strings into integer day counts.

Example: `['1D', '5D'] -> [1, 5]`.

2.6.4 `add_custom_calendar_timedelta(input, timedelta, freq)`

Add a `Timedelta` to a timestamp or `DatetimeIndex` while respecting a custom trading calendar.

Accepted freq values

- `Day`
- `BusinessDay`
- `CustomBusinessDay`

2.6.5 `make_naive_ts(t)`

Return a timezone-naive timestamp.

- If `t` is timezone-aware, it is converted to naive UTC-localized time.
- Otherwise it is localized to `None`.

2.6.6 `diff_custom_calendar_timedeltas(start, end, freq)`

Compute the effective elapsed time between two timestamps under a custom calendar.

This is used when forward-return horizons must respect trading days, weekends, and holidays.

Accepted freq values

- any `pandas.tseries.offsets.BaseOffset`
- commonly `Day`, `BusinessDay`, or `CustomBusinessDay`

2.7 Column detection and display helpers

2.7.1 `get_forward_returns_columns(columns, require_exact_day_multiple=False)`

Identify which columns look like forward-return horizons.

Behavior

- Recognizes `Timedelta`-style labels such as `1D`, `5D`, `30m`, `1D1h`.
- When `require_exact_day_multiple=True`, only exact day multiples are kept.
- Emits a warning if non-day-multiple columns are skipped in that mode.

2.7.2 `print_table(table, name=None, fmt=None)`

Pretty-print a `Series` or `DataFrame`.

Behavior

- Uses rich display output when available.
- Falls back to standard formatted output.
- Temporarily changes `pandas` floating-point formatting when `fmt` is provided.

2.8 Error handling helper

2.8.1 `rethrow(exception, additional_message)`

Re-raise an exception while preserving the original stack trace and appending extra context to the message.

This is used internally to make binning errors easier to diagnose.

2.9 Practical usage

Typical `Alphas` usage looks like this:

1. Prepare a factor series indexed by `(date, asset)`.
2. Prepare a price `DataFrame` with assets in columns.
3. Call `get_clean_factor_and_forward_returns()`.
4. Pass the result into tear sheets or plotting helpers.

Example call: `get_clean_factor_and_forward_returns(factor=..., prices=..., periods=(1, 5, 10))`.

2.10 Notes

- Several functions assume data is already aligned on `(date, asset)`.
- Zero-aware binning is intended for factors centered around zero.
- `filter_zscore` can introduce lookahead bias, so use it carefully.
- Forward-return horizon labels are derived from the price calendar and may reflect custom business-day offsets.

3 `alphalens.performance`

Performance-analysis utilities for factor research, including information coefficient calculations, factor-weighted portfolio simulation, quantile analysis, event-study helpers, and Pyfolio-ready outputs.

This module is the main analytics layer used after factor data has been cleaned and aligned with `alphalens.utils.get_clean_factor_and_forward_returns()`.

3.1 Data model used by this module

Most functions expect a MultiIndex `pandas` object indexed by:

- level 0: `date`
- level 1: `asset`

The input usually includes:

- a `factor` column
- one or more forward-return columns such as 1D, 5D, or 10D
- `factor_quantile`
- optionally `group`

Forward-return column names are detected using `alphalens.utils.get_forward_returns_columns()`.

3.2 Core analysis functions

3.2.1 `factor_information_coefficient(factor_data, group_adjust=False, by_group=False)`

Compute the Spearman rank correlation between factor values and each forward-return horizon.

Behavior

- Calculates the Information Coefficient (IC) period by period.
- Optionally demeanes forward returns by group before computing IC.
- Optionally computes IC separately for each group.
- Preserves the factor data frequency on the date index when not grouped by asset group.

Returns A `DataFrame` of IC values indexed by date, or by date and group when `by_group=True`.

3.2.2 `mean_information_coefficient(factor_data, group_adjust=False, by_group=False, by_time=None)`

Compute the mean IC across the full sample or over a time window.

Behavior

- Delegates to `factor_information_coefficient()`.
- Can aggregate by time period using a `pandas` time rule such as monthly or weekly windows.
- Can aggregate by group when `by_group=True`.

Returns A scalar-like `Series/DataFrame` of mean IC values depending on the requested grouping.

3.3 Portfolio weighting and return simulation

3.3.1 `factor_weights(factor_data, demeaned=True, group_adjust=False, equal_weight=False)`

Compute asset weights from factor values.

Behavior

- Factor values are normalized to gross leverage of 1.
- When `demeaned=True`, weights are constructed as a long-short portfolio.
- When `group_adjust=True`, weights are computed in a group-neutral way.
- When `equal_weight=True`, assets are equal-weighted instead of factor-weighted.

Returns A Series of weights indexed by `(date, asset)`.

Worked example Suppose one date contains three assets with these factor values:

asset	factor
A	2.0
B	1.0
C	-1.0

With the default `demeaned=True`, `group_adjust=False`, and `equal_weight=False`:

1. Demean factor values by date.
 - Mean factor = $(2.0 + 1.0 - 1.0)/3 = 0.6667$
 - Demeaned values = A: 1.3333, B: 0.3333, C: -1.6667
2. Normalize by the sum of absolute values.
 - Gross exposure = $|1.3333| + |0.3333| + |-1.6667| = 3.3333$
 - Weights = A: 0.40, B: 0.10, C: -0.50

These weights represent a dollar-neutral long-short portfolio: long positions (A: 0.40, B: 0.10) total 0.50 and short positions (C: -0.50) total 0.50 in absolute value.

If `demeaned=False`, weights come directly from the raw (non-demeaned) factor values: the sum of absolute raw factors would be $|2.0| + |1.0| + |-1.0| = 4.0$, yielding weights A: 0.50, B: 0.25, C: -0.25. If `group_adjust=True`, weights are adjusted to be group-neutral. If `equal_weight=True`, assets are equal-weighted within long and short buckets instead of using factor magnitudes.

3.3.2 `factor_returns(factor_data, demeaned=True, group_adjust=False, equal_weight=False, by_asset=False)`

Compute period-wise returns for a factor-weighted portfolio.

In this module, a **factor return** is the return of a simulated portfolio whose asset weights are derived from the factor signal. For each date and each forward-return horizon, the return is computed as the sum of:

$$\text{asset weight} \times \text{asset forward return}$$

Equivalently:

$$\text{factor return}[t, \text{period}] = \sum_i w_{t,i} \cdot r_{t,i,\text{period}}$$

Behavior

- Uses `factor_weights()` to build the portfolio.
- Multiplies weights by forward-return columns.
- If `by_asset=True`, returns the weighted return contribution of each asset.
- Otherwise, aggregates to date-level portfolio returns.

Worked example Suppose one date contains three assets with these factor values and 1D forward returns:

asset	factor	1D forward return
A	2.0	0.010
B	1.0	0.020
C	-1.0	-0.010

With the default `demeaned=True`, `group_adjust=False`, and `equal_weight=False`:

1. Demean factor values by date.
 - Mean factor = $(2.0 + 1.0 - 1.0)/3 = 0.6667$
 - Demeaned values = A: 1.3333, B: 0.3333, C: -1.6667
2. Normalize by the sum of absolute values.
 - Gross exposure = $|1.3333| + |0.3333| + |-1.6667| = 3.3333$
 - Weights = A: 0.40, B: 0.10, C: -0.50
3. Multiply each weight by the forward return.
 - A: $0.40 \times 0.010 = 0.0040$
 - B: $0.10 \times 0.020 = 0.0020$
 - C: $-0.50 \times -0.010 = 0.0050$
4. Sum across assets.
 - factor return = $0.0040 + 0.0020 + 0.0050 = 0.0110$

So the factor portfolio return for that date and period is 0.011, or 1.1%.

If `demeaned=False`, the same formula is used, but the weights come directly from the raw factor values rather than demeaned values. If `group_adjust=True`, weights are computed in a group-neutral way. If `equal_weight=True`, the code uses equal-weighted long/short buckets instead of factor-magnitude weights.

Returns A DataFrame of period returns, or asset-level weighted returns when `by_asset=True`.

3.3.3 `factor_alpha_beta(factor_data, returns=None, demeaned=True, group_adjust=False, equal_weight=False)`

Estimate alpha and beta for a factor portfolio using OLS regression.

Behavior

- Uses factor portfolio returns as the dependent variable.
- Uses the universe mean return as the explanatory variable.

- If `returns` is not provided, it is computed using `factor_returns()`.
- Annualizes alpha using a 252-trading-day convention.

Returns A `DataFrame` containing at least annualized alpha and beta by forward-return horizon.

3.3.4 `cumulative_returns(returns)`

Convert simple returns into cumulative returns.

This is a thin wrapper around `empyrical.cum_returns()` with a starting value of 1.

3.3.5 `positions(weights, period, freq=None)`

Build a time series of portfolio positions from factor weights.

Behavior

- Treats weights as active for a holding period defined by `period`.
- Uses the provided trading calendar frequency, or infers one from the weight index.
- Falls back to `BDay` and emits a warning if no frequency is available.
- Recomputes portfolio weights as positions roll forward through time.

Returns A `DataFrame` with timestamps on the index and assets on the columns.

3.4 Quantile and bucket analysis

3.4.1 `mean_return_by_quantile(factor_data, by_date=False, by_group=False, demeaned=True, group_adjust=False)`

Compute mean forward returns and standard errors by factor quantile.

Behavior

- Can compute results by date or across the full sample.
- Can compute results by group.
- Can demean by the whole universe or within each group.

Returns A pair: `(mean_ret, std_error_ret)`.

3.4.2 `compute_mean_returns_spread(mean_returns, upper_quant, lower_quant, std_err=None)`

Compute the difference in mean returns between two quantiles.

Behavior

- Subtracts lower-quantile mean returns from upper-quantile mean returns.
- Optionally propagates standard error for the spread.

Returns A pair: `(mean_return_difference, joint_std_err)`.

3.4.3 `quantile_turnover(quantile_factor, quantile, period=1)`

Measure the proportion of names that leave a given quantile over time.

Behavior

- Compares membership in the selected quantile against the prior period.
- Preserves the date frequency of the input.

Worked example Suppose we evaluate `quantile=5` with `period=1`, and the names in quantile 5 are:

date	names in quantile 5
2026-01-02	$\{A, B, C\}$
2026-01-03	$\{B, C, D\}$
2026-01-06	$\{C, D, E\}$

For each date after the first, turnover is:

$$\frac{\# \text{ new names vs previous date}}{\text{current quantile size}}$$

- On 2026-01-03, new names vs 2026-01-02 are $\{D\}$, so turnover is $1/3 = 0.3333$.
- On 2026-01-06, new names vs 2026-01-03 are $\{E\}$, so turnover is $1/3 = 0.3333$.

The first date has no prior comparison point, so no turnover value is reported for it.

Companion example (`period=2`) Using the same memberships, each date is compared to the set two dates earlier.

- On 2026-01-06, compare $\{C, D, E\}$ to 2026-01-02 ($\{A, B, C\}$).
- New names are $\{D, E\}$, so turnover is $2/3 = 0.6667$.

This highlights that larger `period` values measure non-adjacent membership change. If intermediate dates are missing after frequency alignment (for example via `asfreq`), larger `period` comparisons can be skipped or shifted to different comparable dates.

Returns A Series indexed by date.

3.4.4 `factor_rank_autocorrelation(factor_data, period=1)`

Measure the autocorrelation of factor ranks across periods.

Behavior

- Ranks assets by their factor values within each date, then computes the correlation between each date's rank vector and the rank vector period dates earlier.
- Useful as a turnover/stability diagnostic.

Worked example (period=1) Suppose the factor values for assets A, B, C, D are:

date	factor values (A, B, C, D)	rank vector (A, B, C, D)
2026-01-02	(1, 2, 3, 4)	(1, 2, 3, 4)
2026-01-03	(4, 3, 2, 1)	(4, 3, 2, 1)
2026-01-06	(1, 2, 3, 4)	(1, 2, 3, 4)

With `period=1`, each date is correlated with the immediately prior date:

- On 2026-01-03: $\text{corr}(4, 3, 2, 1), (1, 2, 3, 4) = -1.0$.
- On 2026-01-06: $\text{corr}(1, 2, 3, 4), (4, 3, 2, 1) = -1.0$.

The first date has no prior rank vector, so its value is `NaN`.

Companion example (period=2) Using the same data with `period=2`, 2026-01-06 is compared to 2026-01-02:

- $\text{corr}(1, 2, 3, 4), (1, 2, 3, 4) = 1.0$.

So a larger `period` can reveal longer-horizon rank stability even when adjacent dates are unstable.

Returns A Series of autocorrelation values indexed by date.

3.5 Event-study helpers

3.5.1 `common_start_returns(factor, returns, before, after, cumulative=False, mean_by_date=False, demean_by=None)`

Align return windows around common event dates.

Behavior

- Builds a return window around each factor date and asset pair.
- Aligns all windows to a common event-time index.
- Can work with cumulative or period returns.
- Can de-mean against a reference universe.
- Can average across assets by date.

Returns A DataFrame of aligned return windows.

3.5.2 `average_cumulative_return_by_quantile(factor_data, returns, periods_before=10, periods_after=15, demeaned=True, group_adjust=False, by_group=False)`

Compute average cumulative returns around factor events by quantile.

Behavior

- Uses `common_start_returns()` internally.
- Computes mean and standard deviation of event-time cumulative returns.
- Can separate results by group.
- Supports group-neutral and universe-demeaned variants.

Returns A MultiIndex DataFrame containing mean and standard deviation across the event window.

3.6 Portfolio simulation outputs

3.6.1 `factor_cumulative_returns(factor_data, period, long_short=True, group_neutral=False, equal_weight=False, quantiles=None, groups=None)`

Simulate a factor portfolio and return cumulative performance.

Behavior

- Filters to a single forward-return horizon specified by `period`.
- Can limit analysis to selected quantiles or groups.
- Uses `factor_returns()` and then converts the result to cumulative returns.

This means cumulative factor returns are built directly from the period-wise factor return series described above.

Returns A cumulative return Series.

3.6.2 `factor_positions(factor_data, period, long_short=True, group_neutral=False, equal_weight=False, quantiles=None, groups=None)`

Simulate a factor portfolio and return the time series of positions.

Behavior

- Filters to the requested forward-return horizon.
- Reuses `factor_weights()` to compute holdings.
- Converts weights into rolling positions with `positions()`.

Returns A DataFrame of asset positions over time.

3.6.3 `create_pyfolio_input(factor_data, period, capital=None, long_short=True, group_neutral=False, equal_weight=False, quantiles=None, groups=None, benchmark_period="1D")`

Create returns, positions, and benchmark data in the format expected by Pyfolio.

Behavior

- Builds cumulative strategy returns.
- Resamples returns and positions to daily frequency.
- Adds a `cash` column to the positions output.
- Optionally converts percentage positions into dollar positions using `capital`.
- Computes a benchmark series from the factor universe when the benchmark period is available.

Returns A tuple: (returns, positions, benchmark).

3.7 Dependencies and implementation notes

- Uses `pandas`, `numpy`, `scipy.stats`, `statsmodels`, and `empyrical`.
- Relies heavily on `alphalens.utils` for forward-return column detection and return demeaning.
- Most functions assume data has already been cleaned and aligned into the Alphalens Multi-Index format.
- Several functions use a trading-calendar frequency attached to the date index, so preserving index frequency is important for correct behavior.

3.8 Practical usage

Typical workflow:

1. Build cleaned factor data with `alphalens.utils.get_clean_factor_and_forward_returns()`.
2. Compute diagnostics such as IC, mean returns by quantile, and turnover.
3. Simulate factor-weighted returns or positions.
4. Feed the outputs into plotting or Pyfolio workflows.

Example call: `factor_information_coefficient(factor_data)` or `create_pyfolio_input(factor_data, period="1D")`.

4 alphalens.plotting

Visualization utilities for factor research, including plots for Information Coefficient analysis, factor returns, quantile performance, and other diagnostic charts. The module provides publication-quality visualizations with sane defaults for the Alphalens workflow.

This module visualizes the outputs from `alphalens.performance` and the cleaned factor data from `alphalens.utils`.

4.1 Styling and context management

4.1.1 `customize(func)` — decorator

Apply default plotting context and style to a function.

Behavior

- Sets seaborn colorblind-friendly palette.
- Applies default plotting context and axes style.
- Can be disabled with `set_context=False` in the decorated function call.

4.1.2 `plotting_context(context="notebook", font_scale=1.5, rc=None)`

Create an alphalens-specific plotting context.

Behavior

- Wraps `seaborn.plotting_context()` with alphalens defaults.
- Default line width is 1.5.
- Can be used as a context manager for multiple plots.

Returns A seaborn plotting context object suitable for use in `with` statements.

4.1.3 `axes_style(style="darkgrid", rc=None)`

Create an alphalens-specific axes style context.

Behavior

- Wraps `seaborn.axes_style()` with alphalens defaults.
- Commonly used with `plotting_context()`.

Returns A seaborn axes style object suitable for use in `with` statements.

4.2 Data table displays

4.2.1 `plot_returns_table(alpha_beta, mean_ret_quantile, mean_ret_spread_quantile, return_df=False)`

Display a summary table of factor portfolio returns, alpha, beta, and quantile statistics.

Behavior

- Combines alpha/beta results with quantile return statistics.
- Converts decimal returns to basis points.
- Prints a formatted table or returns as a `DataFrame`.

Returns Pretty-printed table or a `DataFrame` when `return_df=True`.

4.2.2 `plot_turnover_table(autocorrelation_data, quantile_turnover, return_df=False)`

Display quantile turnover and factor rank autocorrelation statistics.

Behavior

- Shows mean turnover for each quantile by period.
- Shows mean factor rank autocorrelation.
- Prints formatted tables or returns as `DataFrames`.

Returns Printed tables or a pair of `DataFrames` when `return_df=True`.

4.2.3 `plot_information_table(ic_data, return_df=False)`

Display Information Coefficient summary statistics.

Behavior

- Computes IC mean, standard deviation, risk-adjusted IC, and t-statistics.
- Includes skewness and kurtosis.
- Prints a summary table or returns as a `DataFrame`.

Returns Pretty-printed table or a `DataFrame` when `return_df=True`.

4.2.4 `plot_quantile_statistics_table(factor_data, return_df=False)`

Display factor value statistics by quantile.

Behavior

- Groups factor data by `factor_quantile`.
- Reports min, max, mean, standard deviation, and count per quantile.
- Includes count percentages.

Returns Pretty-printed table or a `DataFrame` when `return_df=True`.

4.3 Information Coefficient visualization

4.3.1 `plot_ic_ts(ic, ax=None)`

Plot the Information Coefficient time series with a rolling 22-day average.

Behavior

- Creates one subplot per forward-return period.
- Overlays daily IC and a 22-day moving average.
- Includes mean and standard deviation annotations.
- Aligns y-axis ranges across subplots.

Returns A numpy array of matplotlib axes.

4.3.2 `plot_ic_hist(ic, ax=None)`

Plot IC as a histogram with kernel density estimate.

Behavior

- Creates a 3-column grid of histograms, one per forward-return period.
- Includes KDE overlay and a vertical line at the mean.
- Clamps x-axis to `[-1, 1]`.

Returns A numpy array of matplotlib axes.

4.3.3 `plot_ic_qq(ic, theoretical_dist=scipy.stats.norm, ax=None)`

Plot IC values against a theoretical distribution using a Q-Q plot.

Behavior

- Creates a 3-column grid of Q-Q plots.
- Supports normal and Student-t distributions.
- Uses `statsmodels.qqplot()` for plotting.

Returns A numpy array of matplotlib axes.

4.3.4 `plot_ic_by_group(ic_group, ax=None)`

Plot IC by group as a bar chart.

Behavior

- Shows mean IC for each group.
- Rotates x-axis labels for readability.

Returns A matplotlib axes object.

4.3.5 `plot_monthly_ic_heatmap(mean_monthly_ic, ax=None)`

Plot IC or returns as a month/year heatmap with color coding.

Behavior

- Creates one heatmap per forward-return period.
- Uses reversed coolwarm colormap (red = negative, blue = positive).
- Annotates cell values.
- Centered at 0.

Returns A numpy array of matplotlib axes.

4.4 Returns analysis

4.4.1 `plot_quantile_returns_bar(mean_ret_by_q, by_group=False, ylim_percentiles=None, ax=None)`

Plot mean returns by factor quantile as a bar chart.

Behavior

- Optionally creates separate subplots for each group.
- Can limit y-axis range using percentiles of observed data.
- Converts to basis points.

Returns A matplotlib axes object or array depending on grouping.

4.4.2 `plot_quantile_returns_violin(return_by_q, ylim_percentiles=None, ax=None)`

Plot return distributions by factor quantile using violin plots.

Behavior

- Shows distribution of returns within each quantile across periods.
- Uses seaborn violin plot with quartile indicators.
- Optionally limits y-axis range using percentiles.

Returns A matplotlib axes object.

4.4.3 `plot_mean_quantile_returns_spread_time_series(mean_returns_spread, std_err=None, bandwidth=1, ax=None)`

Plot the top-minus-bottom quantile return spread over time.

Behavior

- Shows the difference between top and bottom quantile returns period by period.
- Overlays a 22-day rolling average.
- Optionally adds confidence bands based on standard error.
- Symmetrizes y-axis around zero at the 95th percentile.

Returns A matplotlib axes object or array if input is a DataFrame.

4.4.4 `plot_cumulative_returns(factor_returns, period, freq=None, title=None, ax=None)`

Plot cumulative returns of a factor-weighted portfolio.

Behavior

- Converts simple returns to cumulative using `performance.cumulative_returns()`.
- Includes a reference line at 1.0 (break-even).
- Displays the forward-return period in the title.

Returns A matplotlib axes object.

4.4.5 `plot_cumulative_returns_by_quantile(quantile_returns, period, freq=None, ax=None)`

Plot cumulative returns separately for each factor quantile.

Behavior

- Uses reversed coolwarm colormap for quantile coloring.
- Uses symmetric log scale on the y-axis for readability across wide ranges.
- Shows all quantiles on a single plot with a legend.

Returns A matplotlib axes object.

4.4.6 `plot_quantile_average_cumulative_return(avg_cumulative_returns, by_quantile=False, std_bar=False, title=None, ax=None)`

Plot average cumulative returns by factor quantile around event dates.

Behavior

- Shows mean cumulative returns across the event window for each quantile.
- Optionally creates separate subplots per quantile for clarity.
- Can overlay standard deviation error bars.
- Uses reversed coolwarm colormap for consistency.
- Marks event date ($x=0$) with a vertical dashed line.

Returns A matplotlib axes object or array.

4.5 Factor characteristics

4.5.1 `plot_factor_rank_auto_correlation(factor_autocorrelation, period=1, ax=None)`

Plot factor rank autocorrelation over time.

Behavior

- Shows the stability of relative factor rankings.
- Includes a reference line at 0 (no autocorrelation).
- Annotates the mean autocorrelation in a text box.

Returns A matplotlib axes object.

4.5.2 `plot_top_bottom_quantile_turnover(quantile_turnover, period=1, ax=None)`

Plot turnover for the top and bottom quantiles over time.

Behavior

- Extracts the maximum and minimum quantiles.
- Shows the proportion of names new to each quantile each period.
- Useful for detecting data quality or stationarity issues.

Returns A matplotlib axes object.

4.5.3 `plot_events_distribution(events, num_bars=50, ax=None)`

Plot the distribution of factor events across time.

Behavior

- Divides the time span into `num_bars` intervals.
- Counts events in each interval.
- Displays as a bar chart.

Returns A matplotlib axes object.

4.6 Constants

4.6.1 `DECIMAL_TO_BPS = 10000`

Conversion factor for decimal returns to basis points. Used throughout the module when displaying returns-related plots.

4.7 Dependencies and conventions

- **Visualization:** `seaborn`, `matplotlib`, `scipy.stats`
- **Analysis:** `pandas`, `numpy`, `statsmodels`
- **Internal:** `alphalens.utils`, `alphalens.performance`

Color conventions

- **Coolwarm reversed:** Red represents low/negative quantiles; blue represents high/positive quantiles.
- **Colorblind palette:** All plots use seaborn's colorblind palette for accessibility.
- **Line colors:** Green (forestgreen) for means and averages; blue (steelblue) for confidence bands; red (orangered) for rolling averages.

Plot conventions

- **Multi-subplot layouts:** Plots with multiple periods or groups create grid layouts (usually 1×3 or 2×2).
- **Shared axes:** Where applicable, y-axes are shared across subplots to aid comparison.
- **Reference lines:** Zero lines are included on relevant plots to indicate break-even or neutral positions.
- **Log scales:** Cumulative returns plots use symmetric log scale (`symlog`) to handle very large or very small returns.
- **Annotations:** Statistical summaries (mean, std dev) are included in text boxes on many plots.

4.8 Practical usage

Typical workflow:

1. Compute factor diagnostics with `alphalens.performance` functions.
2. Call plotting functions with the results.
3. All plotting functions are decorated with `@customize` to apply consistent styling.
4. Optional: use `set_context=False` in any plotting call to skip automatic styling.

Example call: `plot_ic_ts(ic_data)` or `plot_quantile_returns_bar(mean_ret_by_q)`.

4.9 Notes

- All plotting functions accept an optional `ax` parameter for custom subplot layouts.
- Return values vary: some functions print tables directly, others return `DataFrames` or axes.
- The `return_df=False` parameter on table functions controls printing vs. returning.
- Most functions are designed for Jupyter notebooks but work in any matplotlib environment.

5 alphalens.tears

Tear sheet generation for comprehensive factor analysis. A tear sheet is a multi-plot report that visualizes factor performance from multiple perspectives, including returns analysis, information coefficient statistics, turnover metrics, and event studies.

This module orchestrates the workflows of `alphalens.performance` and `alphalens.plotting` to produce production-quality analysis reports suitable for presentation and decision-making.

5.1 Helper class

5.1.1 GridFigure

Utility class for managing matplotlib grid-based subplot layouts with automatic row/column navigation.

Methods

- `__init__(rows, cols)` — Create a grid figure with specified dimensions.
- `next_row()` — Get the next full-width subplot (advances to the next row).
- `next_cell()` — Get the next cell in the current row (auto-wraps to next row when full).
- `close()` — Close the figure and clean up resources.

Behavior

- Default figure size is 14 inches wide by 7 inches per row.
- Horizontal and vertical spacing are set to 0.4 and 0.3 respectively for readability.

5.2 Tear sheet generators

All tear sheet functions are decorated with `@plotting.customize` to apply consistent styling automatically. The `set_context=False` parameter can be passed to any tear sheet to skip the decorator.

5.2.1 `create_summary_tear_sheet(factor_data, long_short=True, group_neutral=False)`

Generate a compact tear sheet with key returns, information, and turnover statistics.

Behavior

- Displays factor quantile statistics table.
- Computes and shows alpha, beta, top/bottom quantile returns, and spreads.
- Plots factor quantile returns bar chart.
- Computes and displays IC summary statistics.
- Analyzes turnover and factor rank autocorrelation.
- All metrics respect the `long_short` and `group_neutral` flags.

Returns Displays a matplotlib figure with the complete summary.

5.2.2 `create_returns_tear_sheet(factor_data, long_short=True, group_neutral=False, by_group=False)`

Detailed tear sheet focused on returns analysis by factor quantile.

Behavior

- Computes factor portfolio returns.
- Displays alpha and beta metrics.
- Shows mean quantile returns as a bar chart.
- Plots return distributions using violin plots (by date across periods).
- If '1D' returns are available, plots cumulative portfolio returns and quantile-level cumulative returns.
- Displays top-minus-bottom quantile return spread over time with optional error bands.
- When `by_group=True`, creates separate bar charts for each asset group.

Parameters

- `long_short` — Enable long-short demeaning for dollar-neutral analysis.
- `group_neutral` — Normalize weights and returns across groups.
- `by_group` — Create separate visualizations per group.

Returns Displays one or more matplotlib figures.

5.2.3 `create_information_tear_sheet(factor_data, group_neutral=False, by_group=False)`

Tear sheet focused on Information Coefficient (IC) analysis.

Behavior

- Displays IC summary statistics (mean, std, risk-adjusted IC, t-stat, p-value, skewness, kurtosis).
- Plots IC time series with 22-day rolling average for each forward-return period.
- Plots IC histograms with KDE.
- Plots IC Q-Q plots against normal distribution.
- When `by_group=False` (default), creates a heatmap of monthly mean IC.
- When `by_group=True`, shows mean IC separately for each group.

Parameters

- `group_neutral` — Demean forward returns by group before computing IC.
- `by_group` — Separate analysis by asset group.

Returns Displays matplotlib figure(s).

5.2.4 `create_turnover_tear_sheet(factor_data, turnover_periods=None)`

Tear sheet analyzing factor turnover and rank stability.

Behavior

- Displays quantile turnover statistics (mean turnover per quantile by period).
- Displays factor rank autocorrelation statistics.
- Plots time series of top and bottom quantile turnover for each period.

- Plots factor rank autocorrelation over time for each period.
- Automatically extracts day-multiple forward-return periods if `turnover_periods` is not provided.

Parameters

- `turnover_periods` — Custom periods for turnover analysis. If not provided, uses exact day multiples from `factor_data`.

Returns Displays matplotlib figure with turnover and autocorrelation plots.

5.2.5 `create_full_tear_sheet(factor_data, long_short=True, group_neutral=False, by_group=False)`

Comprehensive tear sheet combining all analysis perspectives.

Behavior

- Displays factor quantile statistics.
- Calls `create_returns_tear_sheet()` with all parameters passed through.
- Calls `create_information_tear_sheet()` with appropriate flags.
- Calls `create_turnover_tear_sheet()` with defaults.
- All component tear sheets skip automatic styling via `set_context=False` to apply consistent styling once.

Parameters

- `long_short` — Enable long-short demeaning.
- `group_neutral` — Enable group-neutral normalization.
- `by_group` — Separate analysis per group.

Returns Displays multiple matplotlib figures for all analyses.

5.2.6 `create_event_returns_tear_sheet(factor_data, returns, avgretplot=(5, 15), long_short=True, group_neutral=False, std_bar=True, by_group=False)`

Tear sheet analyzing average cumulative returns around factor events (event study).

Behavior

- Computes average cumulative returns within a window before and after each factor event.
- Plots overall average cumulative returns with optional error bands.
- When `std_bar=True`, plots separate cumulative return traces for each quantile with standard deviation error bars.
- When `by_group=True`, repeats the analysis separately for each asset group.
- Useful for analyzing the predictive power of a factor over different time horizons relative to the event date.

Parameters

- `factor_data` — MultiIndex factor data (no forward returns required for this tear sheet).
- `returns` — DataFrame of asset returns indexed by date.
- `avgretplot` — Tuple of (`periods_before`, `periods_after`) for the event window.
- `long_short` — Demean cumulative returns.
- `group_neutral` — Normalize returns by group.
- `std_bar` — Include error bar plots by quantile.
- `by_group` — Separate analysis per group.

Returns Displays matplotlib figure(s) showing event-window average returns.

5.2.7 `create_event_study_tear_sheet(factor_data, returns, avgretplot=(5, 15), rate_of_ret=True, n_bars=50)`

Tear sheet for analyzing specific events or signals.

Behavior

- Displays factor quantile statistics.
- Plots histogram of event distribution across time (divided into `n_bars` intervals).
- Calls `create_event_returns_tear_sheet()` internally for average cumulative returns analysis.
- Computes and displays mean quantile returns.
- Plots return distributions via violin plots by quantile.
- Can optionally annualize returns using `rate_of_ret=True`.

Parameters

- `factor_data` — MultiIndex factor data indexed by event dates.
- `returns` — Asset returns DataFrame.
- `avgretplot` — Window for event-study cumulative returns.
- `rate_of_ret` — Convert simple returns to annualized rates.
- `n_bars` — Number of time intervals for event distribution histogram.

Returns Displays matplotlib figures for event analysis.

5.3 Data model

All tear sheet functions expect `factor_data` to be a MultiIndex DataFrame with:

- Index: (`date`, `asset`)
- Columns:
 - Forward-return columns (e.g., 1D, 5D, 10D)
 - `factor` — The factor values
 - `factor_quantile` — Quantile assignments
 - `group` (optional) — Asset group identifiers

This is the standard output format from `alphalens.utils.get_clean_factor_and_forward_returns()`.

5.4 Dependencies and concepts

- **Visualization:** matplotlib, seaborn
- **Analysis:** pandas, numpy, scipy.stats, statsmodels
- **Internal:** alphasens.plotting, alphasens.performance, alphasens.utils

Key tear sheet concepts

- **Long-short analysis:** Factor values are demeaned, creating equal-weight long and short sides.
- **Group-neutral analysis:** Asset groups are weighted equally, and returns are normalized within groups.
- **Quantile analysis:** Assets are divided into quantiles/buckets based on factor values; returns and metrics are computed per quantile.
- **Information Coefficient:** Spearman rank correlation between factor values and forward returns; higher IC indicates stronger predictive power.
- **Turnover:** Proportion of assets entering/leaving a quantile each period; low turnover suggests stable rankings.
- **Event study:** Analysis of returns in a window around specific factor events or signals.

5.5 Practical usage

Typical workflow:

1. Prepare cleaned factor data with `alphasens.utils.get_clean_factor_and_forward_returns()`.
2. Call the desired tear sheet function.
3. For a quick overview, use `create_summary_tear_sheet()`.
4. For detailed analysis, use `create_full_tear_sheet()`.
5. For event-based analysis, use `create_event_returns_tear_sheet()` or `create_event_study_tear_sheet()`.

Example:

```
from alphasens.tears import create_full_tear_sheet

create_full_tear_sheet(factor_data, long_short=True, group_neutral=False)
```

5.6 Notes

- All tear sheets automatically apply consistent styling via the `@plotting.customize` decorator.
- `GridFigure` handles the layout complexity automatically; most workflows don't require direct use.
- The `by_group` parameter enables side-by-side or panel analysis when asset groups are present.
- For large datasets or many forward-return periods, tear sheets can take time to render; `create_summary_tear_sheet()` is faster for quick feedback.
- All plotting occurs within the tear sheet functions; call `plt.show()` or allow the Jupyter kernel to display results automatically.